



6th PACIFIC-RIM SYMPOSIUM ON IMAGE AND VIDEO TECHNOLOGY

October 28 - November 1st. Guanajuato, Mexico. 2013



Tutorial: OpenCV & CUDA

Presented by:

Ramon Aranda, Francisco Hernandez-Lopez, Francisco Madrigal,
{arac, fcoj23, pacomd}@cimat.mx



Centro de Investigación en Matemáticas, A.C.

Guanajuato, Gto. October 2013



Outline



- OpenCV & Cuda (Brief Introduction)..... (15 min)
- Image processing in OpenCV (7.5 min)
- Memory allocation in the GPU..... (7.5 min)
- Memory passing between OpenCV and CUDA..... (10 min)
- Operation on parallel (GPU management)(5 min)
- Operations on GPU: First Examples
 - Addition of Vectors/Matrices.....(20 min)
 - Considerations(10 min)

- Parallel Image processing
 - Compose images(20 min)
 - Gradient magnitude.....(20 min)
 - Image filtering.....(35 min)
 - Corner detector..... (20 min)
 - Diffusion image.....(25 min)
- Native Functions of OpenCV that use CUDA: gpu::mat..(15min)
- Parallel Image processing using multiple GPUs: Examples(20min)
- Conclusions: Potential applications.....(10 min)

Motivation:

Common Tasks on Image Processing

- Image filtering
- Stereo Matching
- Morphology
- HOG
- Segmentation
- Etc.

- All Highly Parallelizable




Motivation: OpenCV & CUDA


- You can solve problems:
 - Finance
 - Image processing and Video
 - Linear Algebra, optimization problems
 - Physics, Chemistry, Biology
 - Etc....



Object detection




Finite element methods



Medical Image Processing

Motivation: GPU (using CUDA) vs multi-core CPU



Frame: 2 (IF 1st 0.09s)

Elapsed: 00:00:00

Grid: 189x5x189=0.18MV

Workspace: 714.42KB (98.75MB/5.37GB GPU)

Use GPU 0: Tesla K20c

Frame: 2 (IF 1st 0.10s)

Elapsed: 00:00:00

Grid: 191x6x191=0.22MV

Workspace: 7.68MB

Use Intel Xeon E5-2670 16C/32T

OpenCV & CUDA.

6



Introduction

OpenCV & CUDA




Introduction: What is OpenCV?

- Library of algorithms released under BSD license.
- Interfaces with C++, C, Python and JAVA.
- Can be compiled on Windows, Linux, Android and Mac.
- Has more than 2500 optimized algorithms.
- Support by a big community of users and developers.
- Multiple uses like visual inspection, robotic, etc.

Introduction: How to install OpenCV

- <http://www.opencv.org/>
- <http://www.cmake.org/>



Introduction: OpenCV modules



General Image Processing



Segmentation



Machine Learning, Image Pyramids
Detection




Transforms




Fitting

Video, Stereo, and 3D




Camera Calibration



Features



Depth Maps



Optical Flow



Inpainting



Tracking

Source: www.itseez.com



Introduction: OpenCV modules

- **Contrib:** Miscellaneous contributions
- **Legacy:** Deprecated code
- **Nonfree:** Algorithms with copyright.
- **GPU:** GPU functions (Can use with another CUDA libs)

Introduction: Parallel Computing

- Running more than one calculation at the same time or "in parallel", using more than one processor.




OpenMP

OpenMPI

Cg,
CUDA,
OpenCL

Introduction: GPU

- Flexible and powerful Processor .
- Handles accuracy of (32/64)-bit in floating point.
- Programmed using high level languages.
- Offers lots of GFLOPS.




From CUDA_C_Programming_Guide.pdf



GeForce GTX TITAN

Introduction: GPU

- Specialized for data parallel computing.
- Uses more transistors to data processing than flow control or data storage.



From CUDA_C_Programming_Guide.pdf



Introduction

CUDA: Compute Unified Device Architecture

- GPGPU technology (General-purpose computing on graphics processing units) that lets you use the C programming language to execute code on the graphic processing unit (GPU).
- Developed by NVIDIA.
- To use this architecture it is required to have a GeForce 8 series (or Quadro equivalent), and more recently GPUs.

Introduction: CUDA Features

- Supports the programming language C/C++, Fortran, Matlab, LabView, etc..
- Unification of hardware and software for parallel computing.
- Supports: Single Instruction, Multiple Data (SIMD).
- Libraries for FFT (Fast Fourier Transform), BLAS (Basic Linear Algebra Subroutines), NPP, TRUSTH, CULA, etc.
- Works internally with OpenGL and DirectX.
- Supports operative systems:
 - Windows, Linux and Mac OS.

Introduction: CUDA-Enabled Graphic Cards



Architectures	Capability
8-200 series	1.0-1.3
FERMI (400 series)	2.0-2.1
KEPLER (600 series)	3.0-3.5

Next Architectures (2014-2015)

MaxWell

Volta

GPU Architectures and Capability

See: http://www.nvidia.com/object/cuda_gpus.html

Introduction: Installing CUDA

- Installing CUDA (<http://developer.nvidia.com/cuda/cuda-downloads>)

CUDA Downloads					
CUDA 5.5 PRODUCTION RELEASE					
Operating System	Distribution	Architecture			Related Documentation
		64-bit	x86	ARMv7	
	Vista, 7, 8 - Notebook	64-bit	32-bit		
	Windows Vista, 7, 8 - Desktop	64-bit	32-bit		Windows Getting Started Guide
	XP - Desktop*	64-bit	32-bit		
Linux	RHEL 6	RPM RUN			
	RHEL 5.5	RUN			
	Fedora 18	RPM RUN			Linux Getting Started Guide
	OpenSUSE 12.2	RPM RUN			RPM / DEB Installation Instructions
	SLES 11 (SP1 & SP2)	RPM RUN			
	Ubuntu 12.04	DEB**	RUN	DEB	RUN Installation Instructions
	Ubuntu 12.10	DEB	RUN	DEB	
	Ubuntu 10.04	RUN			
	Mac OSX 10.7,10.8 & 10.9 *NEW*	PKG			Mac Getting Started Guide



Questions?





Image processing in OpenCV

Image processing in OpenCV

• cv :: Mat

- Basic management of matrices

```
1 // make a 7x7 complex matrix filled with 1+3j.  
Mat M(7,7,CV_32FC2, Scalar(1,3));  
3 // and now turn M to a 100x60  
// 15-channel 8-bit matrix.  
5 // The old content will be deallocated  
M.create(100,60,CV_8UC(15));
```

Image processing in OpenCV

- Class `cv::Mat` is responsible for managing the image
- OpenCV provides functions for reading, showing and saving of images.

```
1 #include <opencv2/core/core.hpp>
2 #include <opencv2/highgui/highgui.hpp>
3
4 int main()
5 {
6
7     // read an image
8     cv::Mat image= cv::imread("img.jpg");
9
10    // create image window named "My Image"
11    cv::namedWindow("My_Image");
12
13    // show the image on window
14    cv::imshow("My_Image", image);
15
16    // wait key for 5000 ms
17    cv::waitKey(5000);
18
19    return 0;
20 }
```

Image processing in OpenCV

Pixel access

- There are different ways to access the pixels within an instance of cv:: Mat.
For example, for grayscale images, we can use the member function
“.at<type>”(row,col)

```
image.at<uchar>(j,i)= value;
```

- In the case of more than one channel

```
image.at<cv::Vec3b>(j,i)[channel]= value;
```



CUDA



Memory allocation in the GPU

Memory allocation in the GPU

- Allocate and free memory
 - **cudaMalloc** ((void**) devPtr, size_t size)
 - **cudaFree** (void *devPtr)
- Those are similar to:
 - Malloc() ..
 - Free() ..

Memory allocation in the GPU

Copy memory.

cudaMemcpy(void *dst, const void *src, size_t count, enum cudaMemcpyKind kind)

Kind:

- cudaMemcpyHostToHost
- cudaMemcpyHostToDevice
- cudaMemcpyDeviceToHost
- cudaMemcpyDeviceToDevice



Memory passing between OpenCV and CUDA


See example in “MemoryManage.cpp”



Operation on parallel (GPU management)

Operation on parallel: Programming Model


- A program that is compiled to run on a graphics card is called the *Kernel*.
- The set of threads that execute a kernel is organized as a **grid** of thread blocks.
- A thread block is a set of threads that can cooperate together:
 - Easy access to shared memory.
 - Synchronously.
 - With a thread identifier ID.
 - Blocks can be arranged for 1, 2 or 3 dimensions.
- A **grid of thread blocks**:
 - It has a limited number of threads in a block.
 - The blocks are identified by an ID.
 - Arrangements can be of 1 or 2 dimensions.



Operation on parallel: Programming Model

- Running on the Host and Device.

Host = CPU
Device = GPU
Kernel = Set of instructions than runs in the device



Operation on parallel: Qualifiers for a kernel

● device

- Runs on the device.
- Called only from the device.

● global

- Runs on the device
- Called only from the host.

Operation on parallel: Qualifiers for variables

- **__device__**
 - Resides in global memory space.
 - Has the lifetime of an application.
 - Lives accessible from all threads within the grid, and from the host through the library at runtime.
- **Others:**
 - **__constant__** (Optionally used with **__device__**)
 - Resides in constant memory space.
 - Has the lifetime of an application.
 - Lives accessible from all threads within the grid, and from the host through the library at runtime.
 - **__shared__** (Optionally used with **__device__**)
 - Lives in shared memory space of a thread block.
 - Has the lifetime of a block.
 - Only accessible from the threads that are within the block.

Operation on parallel: Kernel function calls

- Example function
 - Kernel in the Device:
 - `__global__ void NameFunc(float *parameter, ...);`
 - it must be called as follows:
 - `NameFunc <<< Dg, Db, Ns, St >>> (parameter1,...);`
- **Dg**: Type *dim3*, dimension and size of the grid.
- **Db**: Type *dim3*, dimension and size of each block.
- **Ns**: Type *size_t*, number of bytes in shared memory.
- **St**: Type *cudaStream_t* that indicates which stream will use the kernel.
(Ns and St are optional).



Operation on parallel: Automatically Defined Variables

- All `__global__` and `__device__` functions have access to the following variables:
 - **gridDim** (dim3), indicates the dimension of the grid.
 - **blockIdx** (uint3), indicates the index of the bloque within the grid.
 - **blockDim** (dim3), indicates the dimension of the block.
 - **threadIdx** (uint3), indicates the index of the thread within the block.



Operations on GPU: First Examples

Operations on GPU:

Add One

CPU C

```
void add_one_cpu(float *vector, int N)
{
    int i;
    for (i=0;i<N;i++) {
        vector [j] += 1.0f;
    }
}

void main() {
    ....
    add_one_cpu (a,N);
}
```

CUDA C

```
__global__ void add_one_gpu(float *d_vector, int N)
{
    int i=blockIdx.x*blockDim.x+threadIdx.x;
    if(i < N )
        d_vector[i] += 1.0f;
}
```


void main()

```
dim3 dimBlock(blocksize, 1, 1);
dim3 dimGrid(N/dimBlock.x, 1,1);
add_matrix_gpu<<<dimGrid, dimBlock>>>(a, N);
```

Operations on GPU:

Add One

- Every element in the vector is processing by every thread in each block



Operations On GPU: add vectors

- Add two vectors
 - Create host memory: “a_h”, “b_h” and “c_h”
 - Initialize the vectors “a_h” and “b_h”.
 - Create device memory: “a_d”, “b_d” and “c_d”.
 - Copy memory from host to device of vectors a and b.
 - Add vectors a_d and b_d; the result is saved in vector c_d.
 - Copy memory from device to host of vector c.
 - Finally, show the result.
- See “add_vectors.cpp”


Operations on GPU: add Matrices

- Exercise: The code in “add_matrices.cpp” is incomplete; find and correct the mistake.
- Remember:
 - Create host memory: “a_h”, “b_h” and “c_h”.
 - Initialize “a_h” and “b_h”.
 - Create device memory: “a_d”, “b_d” y “c_d”.
 - Copy memory from host to device.
 - Add matrix in the device.
 - Copy memory from device to host.
 - Finally, show the result.

Operations On GPU:

add Matrices

1,1	1,2	1,3
2,1	2,2	2,3
3,1	3,2	3,3



1	2	3
4	5	6
7	8	9

Indexes in Matrix form

Indexes in Vector form

The formula in C/C++ is

Index_vector = i * #cols + j

1	2	3
4	5	6
7	8	9



1	2	3	4	5	6	7	8	9
---	---	---	---	---	---	---	---	---

Operations on GPU: Considerations

- There are some technique to improve the performance of algorithms on GPU.
- Multiple Data, Single Instruction:
 - 32 threads (warp)
 - Avoid use “if”.
 - Also, avoid “for” with different stop criteria in each thread

```
if()          ← only 2
thread
....  

else          ← 30 threads
...
This takes 2 times!
```



Parallel Image processing

PSW Parallel Image processing: Exercise: Image Composition

- Load two images and reserve memory to the output image.
- Create memory on Device (for the 3 images).
- Copy memory of the Host to Device.
- Loop:
 - Kernel (CUDA_Compose_Images)
 - Return the result on the Host
 - Show the result
- Free the memory

PSW Parallel Image processing: EXERCISE: Gradient Magnitude

- Load the original image in host memory.
- Create device memory: Imag_dev, ImagDx_dev, ImagDy_dev, ImagMG_dev.
- Copy the original image from host to device memory.
- Calculate Dx, Dy and GM in the device.

$$D_x(x, y) = I(x, y) - I(x - 1, y)$$

$$D_y(x, y) = I(x, y) - I(x, y - 1)$$

$$GM(x, y) = \sqrt{D_x^2(x, y) + D_y^2(x, y)}$$


- Copy the result from device to host memory.
- Show the result.

PSW Parallel Image processing: Image filtering

- Example: Mean filter
 - Load the original image in host memory.
 - Create device memory.
 - Copy the original image from host to device memory.
 - Calculate the mean filter.
 - Copy the result from device to host memory.
 - Show the result.

Parallel Image processing: Image filtering

- Mean filter with window size of 3x3:



PSW Parallel Image processing: Image filtering

- Exercises: Gaussian and Laplacian filters
 - Load the original image in host memory.
 - Create device memory.
 - Copy the original image from host to device memory.
 - Calculate the Gaussian or Laplacian filter.
 - Copy the result from device to host memory.
 - Show the result

Gaussian Filter:

$$\begin{pmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{pmatrix}$$


Laplacian Filter:

$$\begin{pmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{pmatrix}$$

Parallel Image processing: Corner detector

- Exercise: Corner detector with the structure tensor

$$\begin{bmatrix} D_x^2 & D_x D_y \\ D_x D_y & D_y^2 \end{bmatrix}$$



PSW Parallel Image processing: Exercise - Diffusion image

- ⊕ Given an image $g(x)$ with noise.
- ⊕ Smooth the image $g(x)$ with the following functional:

$$U[f(x)] = \frac{1}{2} \sum_x [f(x) - g(x)]^2 + \frac{\lambda}{2} \sum_{\langle x,y \rangle} [f(x) - f(y)]^2$$

- ⊕ Differentiating and equating to zero, we obtain:



$$f^{k+1}(x) = \frac{g(x) + \lambda \sum_{y \in N_x} f^k(y)}{1 + \lambda |N_x|}$$

- ⊕ We can solve by:
 - ⊕ Jacobi
 - ⊕ Gauss-Seidel

$|N_x| = \#$ neighborhoods
of pixel x

$$f^0(x) = g(x)$$

PSW Parallel Image processing: Exercise - Diffusion image



PSIVT Parallel Image processing using multiple GPUs: Examples

- GPUs can be controlled by:
 - A single CPU thread
 - Multiple CPU threads



PSIV Parallel Image processing using multiple GPUs: Examples

- Asynchronous calls (kernels, memcopies) don't block switching the GPU.
- The following code will have both GPUs executing concurrently:
 - `cudaSetDevice(0);`
 - `kernel<<<...>>>(...);`
 - `cudaSetDevice(1);`
 - `kernel<<<...>>>(...);`

PSW Parallel Image processing using multiple GPUs: Examples



GPU module design considerations

- Key ideas
 - ⦿ Explicit control of data transfers between CPU and GPU
 - ⦿ Minimization of the data transfers
 - ⦿ Completeness
 - ⦿ Port everything even functions with little speed-up
- ⦿ Solution
 - ⦿ Container for GPU memory with upload/download functionality
 - ⦿ GPU module function take the container as input/output parameters

GPU module design considerations

- Class GpuMat –for storing 2D (**pitched**) data on GPU
 - Interface similar to `cv::Mat()`, supports reference counting
 - Its data is not continuous, extra padding in the end of each row
 - It contains:
 - **data** - Pointer data beginning in GPU memory
 - **step** – Distance in bytes is between two consecutive rows
 - **cols, rows** - Fields that contain image size
 - **upload/download** – Up/down memory from device

OpenCV GPU Module Example

```
Mat frame;  
VideoCapture capture(camera);  
cv::HOGDescriptor hog;  
hog.setSVMDetector(cv::HOGDescriptor::  
    getDefaultPeopleDetector());  
  
capture >> frame;  
  
vector<Rect> found;  
hog.detectMultiScale(frame, found,  
    1.4, Size(8, 8), Size(0, 0), 1.05, 8);
```

Designed very similar!

```
Mat frame;  
VideoCapture capture(camera);  
cv::gpu::HOGDescriptor hog;  
hog.setSVMDetector(cv::HOGDescriptor::  
    getDefaultPeopleDetector());  
  
capture >> frame;  
  
GpuMat gpu_frame;  
gpu_frame.upload(frame);  
  
vector<Rect> found;  
hog.detectMultiScale(gpu_frame, found,  
    1.4, Size(8, 8), Size(0, 0), 1.05, 8);
```



Conclusions:

- CPU

- Incremental improvements (memory caches and complex architectures)
- Few Multi-core (4/8/16)

- GPU

- Highly parallel with 100s of simple cores
- Easier to extend by adding more GPUs
- Continue to grow exponentially!
- Most of the GPUs are cheap!

Conclusions:


- We presented a small introduction of the parallel processing using GPUs.
- There are many sofistecated strategies for make up your GPU-code faster.
- Most problems can be parallelized and are suitable to be run on GPUs
- One has to consider the properties of the GPU (shared memory, cache, compute capability) when designing the kernels

Conclusions: Potential applications



Tracking

Conclusions: Potential applications





Tracking

Conclusions: Potential applications



VScreen

Conclusions: Potential applications



A

A

Tract Estimations from the callosum corpus

Tractography



Questions?





Thank you!